

Companies Module Extraction Guidelines

1. Primary objective

The Companies module should contain **companies where AI is a meaningful part of the company's product, technology, or core business.**

We are **not building a generic global company database.**

Prioritize:

- AI-native companies
- AI model companies
- Generative AI companies
- Agentic AI companies
- AI infrastructure companies
- AI developer/tooling companies
- Computer vision companies
- Voice AI companies
- Robotics/physical AI companies
- AI cybersecurity
- AI healthcare
- AI fintech
- AI enterprise software
- AI research companies
- AI hardware/chips specifically built for AI
- AI platforms and AI-as-a-service companies

A company should **not be included merely because it uses AI internally.**

For example, a normal SaaS, consulting, e-commerce, manufacturing, or fintech company that has added an AI feature should generally **not** be treated as an AI company.

This distinction is important. Tracxn, for example, defines **AI-Native** as companies where AI is the core value proposition and primary product infrastructure, rather than AI simply being added to an existing product. (help.tracxn.com)

2. Source priority — VERY IMPORTANT

Do **not** start with generic company databases and scrape everything.

Tier 1 — AI-specific sources

Start with sources that are already focused specifically on AI companies.

Primary source should be:

There's An AI For That (TAAFT) — use its company ecosystem as one of the major discovery sources.

Then use other AI-specific directories, AI startup lists, AI infrastructure lists, AI funding lists, AI accelerator portfolios, etc.

The goal is:

First build the 8,000 strongest/relevant AI companies from AI-focused sources.

Only after this should we expand the dataset.

Tier 2 — AI-specific sections of major databases

Use:

- **Crunchbase — Artificial Intelligence companies**
- **Crunchbase — Artificial Intelligence startups**
- AI-specific Crunchbase hubs and categories

Crunchbase currently has dedicated AI company/startup datasets of roughly **10K organizations**, making it a particularly useful secondary source for reaching the initial target. ([Crunchbase](#))

Use the **AI-specific datasets**, not the general Crunchbase company database.

Tier 3 — Specialized startup databases

Use sources such as:

- Tracxn
- AI startup databases
- AI accelerator portfolios
- AI VC portfolios
- AI industry reports
- AI infrastructure lists
- AI awards/lists
- AI-specific research/company rankings

Tracxn is particularly useful because it has explicit AI categories such as **Artificial Intelligence, AI Infrastructure, Generative AI, Agentic AI, NLP, AI-native, Computer Vision, Voice AI and AI-as-a-Service**. ([Tracxn](https://www.tracxn.com))

Its AI Infrastructure lists can also be used for discovery. ([w.tracxn.com](https://www.tracxn.com))

3. Generic databases are NOT primary discovery sources

Do **not** use:

- Wikipedia/Wikidata
- Generic company databases
- Generic business directories
- Random company lists
- Generic web directories

as the primary source for discovering companies.

These can be used **later for enrichment and verification**.

For example:

AI-specific source → discover company
Official website → verify/enrich
Crunchbase/Tracxn → funding/valuation/investors
Wikidata → additional metadata
LinkedIn → employee/company information
News → recent developments

This keeps the **source of truth AI-focused** rather than allowing a generic database to determine what qualifies as an AI company.

4. Extraction priority

Build the database in this order:

Phase 1 — Top AI companies

Start with the most important companies globally:

- OpenAI
- Anthropic
- Google DeepMind
- xAI
- NVIDIA
- Mistral
- Cohere
- Perplexity
- DeepSeek
- Scale AI
- Hugging Face
- Cursor/Anysphere
- Databricks
- CoreWeave
- etc.

Don't focus only on valuation. Include **important AI companies even if they are smaller or recently founded**.

Phase 2 — 8,000 AI companies

Our first major target should be approximately:

8,000 genuinely relevant AI companies.

Use the combination of:

TAAFT + Crunchbase AI + Tracxn AI + specialized AI sources + AI VC/accelerator lists

to reach this number.

Crunchbase alone has nearly 10,000 organizations in its AI category, so there is already enough dedicated AI-company coverage to build a strong initial dataset. ([Crunchbase](#))

Phase 3 — Expansion

Only after the first ~8K are properly extracted, verified and deduplicated should we start expanding toward the larger company database.

The priority should be:

8K high-quality AI companies > 50K generic companies.

5. Qualification rule

Before adding a company, the extractor should answer:

"Is AI genuinely central to this company's product/business?"

YES → include

If the answer is clearly yes.

NO → don't include

If AI is merely:

- a feature
 - a marketing term
 - an internal technology
 - a minor product capability
 - something the company recently added
 - unrelated to its primary business
-

6. Important distinction: AI Native vs AI-enabled

Create an internal classification:

AI Native

AI is the company's **core product/business**.

AI Enabled

AI is important but not the company's fundamental product.

AI Adjacent

Company operates in an ecosystem related to AI but isn't primarily an AI company.

For the initial **8K**, prioritize **AI Native** and strong **AI-focused** companies.

This will also make the company profile much more useful. Tracxn's current data model similarly distinguishes AI-Native companies from companies that merely add AI capabilities.

(help.tracxn.com)

7. Data to extract for each company

At minimum:

- Company name
- Logo
- Website
- Country
- City
- Founded year
- Description
- AI category
- AI Native / AI Enabled
- Sector
- Employees
- Funding raised
- Latest funding round
- Valuation
- Investors
- Models
- Tools/products

- Repositories
- News
- Videos
- Fundraises
- Social links

Your screenshots already show a good direction here: **company** → **tools** → **models** → **repositories** → **news** → **videos** → **fundraises** → **investments**.

The company page should ultimately become the **central entity connecting the rest of AIOrbit's ecosystem**.

8. Deduplication is critical

The same company will appear across multiple sources.

For example:

Anysphere
Anysphere, Inc.
Cursor / Anysphere
Anysphere AI

should not become four companies.

Use:

company domain + legal name + aliases

as the primary deduplication mechanism.

The **website domain should be the strongest identifier** wherever possible.

9. Source hierarchy for individual fields

Don't assume one source is correct for everything.

Use:

Company website

→ identity, description, products, AI focus

Crunchbase / Tracxn

→ funding, investors, valuation, company stage

Official announcements

→ funding rounds, acquisitions, launches

LinkedIn

→ employees/location/company information

News

→ recent events

Wikidata/Wikipedia

→ supplementary metadata only

This is particularly important for financial fields. Don't copy a valuation from an arbitrary article when an authoritative company announcement or reliable financial database is available.

10. Final rule for the team

Do not optimize for the number of companies extracted. Optimize for the number of relevant AI companies extracted.

The target should be:

First 8,000 high-quality AI companies → verify → deduplicate → enrich → then expand.

And **do not fill the database with Microsoft, Adobe, telecom companies, banks, manufacturers, etc. simply because they have AI products** unless they genuinely qualify under our AI-company criteria.

That distinction is what will make AIOrbit's Companies module feel like an **AI intelligence platform**, rather than another generic company directory.